

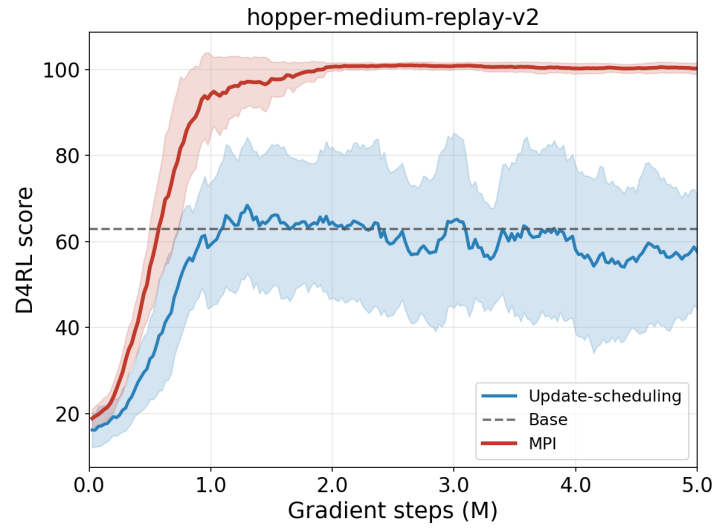


Figure 1: TD3+BC: MPI vs. actor-update scheduling on `hopper-medium-replay-v2`. Base is original TD3+BC; Actor-update scheduling uses five TD3+BC actor updates per delayed actor slot; MPI uses one base update plus four re-centered refinements ($K = 4$). All runs use the original TD3+BC configuration; shaded regions denote standard deviation over 4 seeds.